

## Group Number 215: Black Friday Sales Prediction

First Name	Last Name	Email (hawk.iit.edu)	Student ID
Siddhi	Kulkarni	skulkarni13@hawk.iit.edu	A20430008
Pooja	Choudhari	cpooja@hawk.iit.edu	A20425832
Sagar	Ippili	sippili@hawk.iit.edu	A20417999

### 1. Import data set to R.

```
> data <- read.csv("BlackFriday.csv", sep=";", header=T)
> data[1:5,]
  User_ID Product_ID Gender Age Occupation City_Category Stay_In_Current_City_Years
1 1000001 P00069042      F 0-17          10              A                      2
2 1000001 P00248942      F 0-17          10              A                      2
3 1000001 P00087842      F 0-17          10              A                      2
4 1000001 P00085442      F 0-17          10              A                      2
5 1000002 P00285442      M 55+          16              C                      4+
  Marital_Status Product_Category_1 Product_Category_2 Product_Category_3 Purchase
1              0                  3                  NA              NA      8370
2              0                  1                  6              NA     15200
3              0                 12                  NA              NA      1422
4              0                 12                 14              NA      1057
5              0                  8                  NA              NA      7969
> |
```

### 2. Replace the missing values in columns,

- a. **Product\_Category\_2**: Replacing the missing values with mean value of the column, i.e **9.84**.

```
> summary(data$Product_Category_2)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
  2.00   5.00   9.00   9.84  15.00  18.00 166986
> data$Product_Category_2[is.na(data$Product_Category_2)] <- 9.84
> summary(data$Product_Category_2)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  2.000   8.000   9.840   9.841  14.000  18.000
~ |
```

- b. **Product\_Category\_3**: Replacing the missing values with mean value of the column, i.e **12.7**.

```
> summary(data$Product_Category_3)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
   3.0    9.0   14.0   12.7   16.0   18.0 373299
> data$Product_Category_3[is.na(data$Product_Category_3)] <- 12.7
> summary(data$Product_Category_3)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   3.00  12.70  12.70  12.69  12.70  18.00
> |
```

### 3. Perform Z Test for

- Hypothesis 1:

```
> z.test(data$Product_Category_1,NULL,alternative = "less",mu = 5,sigma.x = sd(data$Product_Category_1),sigma.y = NULL,conf.level = 0.95)
```

One-sample z-Test

```
data: data$Product_Category_1
z = 57.774, p-value = 1
alternative hypothesis: true mean is less than 5
95 percent confidence interval:
NA 5.303961
sample estimates:
mean of x
5.295546
```

- Hypothesis 2:

```
> z.test(data$Product_Category_1,data$Product_Category_2,alternative="two.sided",
+ mu=0,sigma.x=sd(data$Product_Category_1),sigma.y=sd(data$Product_Category_2),conf.level=0.95)
```

Two-sample z-Test

```
data: data$Product_Category_1 and data$Product_Category_2
z = -590.05, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
-4.561032 -4.530831
sample estimates:
mean of x mean of y
5.295546 9.841478
```

4. Ignoring the columns 'User\_ID' and 'Product\_ID', because they are unique to themselves and have no impact on our dependent variable '**Purchase**', followed by creating a new data set for further building classification models.

```
> data<-subset(data,select = -User_ID)
> data<-subset(data,select = -Product_ID)
> data[1:3,]
  Gender Age Occupation City_Category Stay_In_Current_City_Years Marital_Status
1     F 0-17         10             A                           2              0
2     F 0-17         10             A                           2              0
3     F 0-17         10             A                           2              0
  Product_Category_1 Product_Category_2 Product_Category_3 Purchase
1                 3          9.84          12.7      8370
2                 1          6.00          14.0     15200
3                12          9.84          12.7     1422
> |
```

5. Building Naïve-Bayes Classification Model:

- a. Clone the data set to a new variable, 'data\_nb' to use for building Naïve-Bayes classification model.

```

> data_nb<-data
> data_nb[1:5,]
  Gender Age Occupation City_Category Stay_In_Current_City_Years Marital_Status
1      F 0-17          10              A                        2              0
2      F 0-17          10              A                        2              0
3      F 0-17          10              A                        2              0
4      F 0-17          10              A                        2              0
5      M 55+          16              C                        4+              0
  Product_Category_1 Product_Category_2 Product_Category_3 Purchase
1                   3                   9.84                12.7    8370
2                   1                   6.00                14.0   15200
3                   12                   9.84                12.7    1422
4                   12                   14.00               12.7    1057
5                   8                    9.84                12.7   7969
> |

```

- b. Transform the '**Purchase**' variable into a dummy variable which has two values, i.e **1** and **0**. Here, **1** means higher purchase [i.e >Mean Value] amount and **0** means lower purchase [i.e <Mean Value] amount made by a customer.

```

> summary(data_nb$Purchase)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  185    5866    8062    9334   12073   23961
> data_nb$Purchase<-ifelse(data_nb$Purchase>9334,1,0)
> data_nb$Purchase<-as.factor(data_nb$Purchase)
> head(data_nb$Purchase)
[1] 0 1 0 0 0 1
Levels: 0 1
> |

```

- c. Group the numerical variables using the **cut()**, i.e

I. 'Product\_Category\_1':

```

> summary(data_nb$Product_Category_1)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 1.000  1.000   5.000   5.296  8.000  18.000
> data_nb$Product_Category_1<-cut(data_nb$Product_Category_1,3)
> head(data_nb$Product_Category_1)
[1] (0.983,6.67] (0.983,6.67] (6.67,12.3] (6.67,12.3] (6.67,12.3] (0.983,6.67]
Levels: (0.983,6.67] (6.67,12.3] (12.3,18]
> |

```

II. 'Product\_Category\_2':

```

> summary(data_nb$Product_Category_2)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 2.000  8.000   9.840   9.841  14.000  18.000
> data_nb$Product_Category_2<-cut(data_nb$Product_Category_2,4)
> head(data_nb$Product_Category_2)
[1] (6,10] (1.98,6] (6,10] (10,14] (6,10] (1.98,6]
Levels: (1.98,6] (6,10] (10,14] (14,18]
> |

```

III. 'Product\_Category\_3':

```

> summary(data_nb$Product_Category_3)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   3.00  12.70   12.70   12.69  12.70   18.00
> data_nb$Product_Category_3<-cut(data_nb$Product_Category_3,4)
> head(data_nb$Product_Category_3)
[1] (10.5,14.2] (10.5,14.2] (10.5,14.2] (10.5,14.2] (10.5,14.2] (10.5,14.2]
Levels: (2.98,6.75] (6.75,10.5] (10.5,14.2] (14.2,18]
> |

```

- d. Check the entire data set after grouping

```

> head(data_nb)
  Gender  Age Occupation City_Category Stay_In_Current_City_Years Marital_Status
1     F  0-17         10              A                        2              0
2     F  0-17         10              A                        2              0
3     F  0-17         10              A                        2              0
4     F  0-17         10              A                        2              0
5     M   55+         16              C                       4+              0
6     M 26-35         15              A                        3              0
  Product_Category_1 Product_Category_2 Product_Category_3 Purchase
1      (0.983,6.67]      (6,10]      (10.5,14.2]      0
2      (0.983,6.67]      (1.98,6]      (10.5,14.2]      1
3      (6.67,12.3]      (6,10]      (10.5,14.2]      0
4      (6.67,12.3]      (10,14]      (10.5,14.2]      0
5      (6.67,12.3]      (6,10]      (10.5,14.2]      0
6      (0.983,6.67]      (1.98,6]      (10.5,14.2]      1
> |

```

- e. Dataset being very large (approx. 500 thousand records), we will split data into train data and test data for hold-out evaluation later.

```

> index.data_nb<-sample(1:nrow(data_nb),size=round(0.8*nrow(data_nb)))
> train.data_nb<-data_nb[index.data_nb,]
> test.data_nb<-data_nb[-index.data_nb,]
> head(train.data_nb)
  Gender Age Occupation City_Category Stay_In_Current_City_Years Marital_Status
170762  M 36-45         2             A                        2              0
11935   M 26-35         0             C                        1              0
305395  M 46-50        12             C                        1              1
487998  M 36-45        17             C                        3              0
159734  M 18-25         0             B                        1              0
390377  F 46-50        16             A                        1              1
  Product_Category_1 Product_Category_2 Product_Category_3 Purchase
170762      (6.67,12.3]      (10,14]      (10.5,14.2]      0
11935      (0.983,6.67]      (14,18]      (10.5,14.2]      1
305395      (6.67,12.3]      (6,10]      (10.5,14.2]      0
487998      (6.67,12.3]      (6,10]      (10.5,14.2]      0
159734      (0.983,6.67]      (6,10]      (10.5,14.2]      1
390377      (0.983,6.67]      (10,14]      (10.5,14.2]      0
> head(tes.data_nb)
Error in head(tes.data_nb) : object 'tes.data_nb' not found
> head(test.data_nb)
  Gender Age Occupation City_Category Stay_In_Current_City_Years Marital_Status
2      F 0-17         10             A                        2              0
4      F 0-17         10             A                        2              0
19     M 36-45         1             B                        1              1
21     M 26-35        12             C                       4+              1
28     M 26-35        17             C                        0              0
30     F 36-45         1             B                       4+              1
  Product_Category_1 Product_Category_2 Product_Category_3 Purchase
2      (0.983,6.67]      (1.98,6]      (10.5,14.2]      1
4      (6.67,12.3]      (10,14]      (10.5,14.2]      0
19     (0.983,6.67]      (10,14]      (14.2,18]      1
21     (0.983,6.67]      (10,14]      (10.5,14.2]      0
28     (0.983,6.67]      (10,14]      (10.5,14.2]      0
30     (0.983,6.67]      (1.98,6]      (6.75,10.5]      1
> |

```

f. Build the Naïve-Bayes classification model.

```

> model_nb<-naive_bayes(Purchase~.,train.data_nb)
> pred_nb<-predict(model_nb,test.data_nb)
> head(predict(model_nb,test.data_nb,type="prob"))
              0              1
[1,] 0.4773295 0.52267053
[2,] 0.9418721 0.05812791
[3,] 0.5076089 0.49239106
[4,] 0.7416757 0.25832425
[5,] 0.7383431 0.26165689
[6,] 0.1365317 0.86346833
> library(Metrics)
> accuracy(test.data_nb$Purchase,pred_nb)
[1] 0.7089057
> |

```

From the Naïve-Bayes model, we can see that the accuracy is **0.7089057**.

- g. Try building different Naïve-Bayes classification models by categorizing the numerical variables in to different number of groups using the 'cut()'.

- I. Product\_Category\_1 with 3, Product\_Category\_2 with 2 and Product\_Category\_3 with 4 groups each.

```
> data_nb<-data
> data_nb<-subset(data,select = -User_ID)
> data_nb<-subset(data_nb,select = -Product_ID)
> data_nb$Purchase<-ifelse(data_nb$Purchase>9334,1,0)
> data_nb$Purchase<-as.factor(data_nb$Purchase)
> head(data_nb)
  Gender Age Occupation City_Category Stay_In_Current_City_Years Marital_Status
1     F  0-17          10              A                        2              0
2     F  0-17          10              A                        2              0
3     F  0-17          10              A                        2              0
4     F  0-17          10              A                        2              0
5     M   55+          16              C                       4+              0
6     M 26-35          15              A                        3              0
> data_nb$Occupation<-as.factor(data_nb$Occupation)
> data_nb$Marital_Status<-as.factor(data_nb$Marital_Status)
> data_nb$Product_Category_1<-cut(data_nb$Product_Category_1,3)
> data_nb$Product_Category_2<-cut(data_nb$Product_Category_2,2)
> data_nb$Product_Category_3<-cut(data_nb$Product_Category_3,4)
> index.data_nb<-sample(1:nrow(data_nb),size=round(0.8*nrow(data_nb)))
> train.data_nb<-data_nb[index.data_nb,]
> test.data_nb<-data_nb[-index.data_nb,]
> model_nb<-naive_bayes(Purchase~.,train.data_nb)
> pred<-predict(model_nb,test.data_nb)
> head(predict(model_nb,test.data_nb,type="prob"))
      0      1
[1,] 0.2617788 0.7382212
[2,] 0.6622482 0.3377518
[3,] 0.8389146 0.1610854
[4,] 0.7349300 0.2650700
[5,] 0.3308822 0.6691178
[6,] 0.7090515 0.2909485
> accuracy(test.data_nb$Purchase,pred)
[1] 0.6986467
```

From the above Naïve-Bayes model, we can see that the accuracy is 0.6986467.

- II. Product\_Category\_1 with 6, Product\_Category\_2 with 6 and Product\_Category\_3 with 2 groups each.

```
> data_nb<-data_nb_new
> data_nb$Product_Category_2<-cut(data_nb$Product_Category_2,6)
> data_nb$Product_Category_3<-cut(data_nb$Product_Category_3,6)
> data_nb$Product_Category_1<-cut(data_nb$Product_Category_1,2)
> index.data_nb<-sample(1:nrow(data_nb),size=round(0.8*nrow(data_nb)))
> train.data_nb<-data_nb[index.data_nb,]
> test.data_nb<-data_nb[-index.data_nb,]
> model_nb<-naive_bayes(Purchase~.,train.data_nb)
> pred<-predict(model_nb,test.data_nb)
> accuracy(test.data_nb$Purchase,pred)
[1] 0.7030089
```

From the above Naïve-Bayes model, we can see that the accuracy is 0.7030089.

**III. Product\_Category\_1 with 8, Product\_Category\_2 with 8 and Product\_Category\_3 with 6 groups each.**

```
> data_nb<-data_nb_new
> data_nb$Product_Category_3<-cut(data_nb$Product_Category_3,8)
> data_nb$Product_Category_2<-cut(data_nb$Product_Category_2,8)
> data_nb$Product_Category_1<-cut(data_nb$Product_Category_1,6)
> index.data_nb<-sample(1:nrow(data_nb),size=round(0.8*nrow(data_nb)))
> train.data_nb<-data_nb[index.data_nb,]
> test.data_nb<-data_nb[-index.data_nb,]
> model_nb<-naive_bayes(Purchase~.,train.data_nb)
> pred<-predict(model_nb,test.data_nb)
> accuracy(test.data_nb$Purchase,pred)
[1] 0.7899735
```

**From the above Naïve-Bayes model, we can see that the accuracy is 0.7899735.**

**IV. Product\_Category\_1 with 8, Product\_Category\_2 with 8 and Product\_Category\_3 with 8 groups each.**

```
> data_nb<-data_nb_new
> data_nb$Product_Category_3<-cut(data_nb$Product_Category_3,8)
> data_nb$Product_Category_2<-cut(data_nb$Product_Category_2,8)
> data_nb$Product_Category_1<-cut(data_nb$Product_Category_1,8)
> index.data_nb<-sample(1:nrow(data_nb),size=round(0.8*nrow(data_nb)))
> train.data_nb<-data_nb[index.data_nb,]
> test.data_nb<-data_nb[-index.data_nb,]
> model_nb<-naive_bayes(Purchase~.,train.data_nb)
> pred<-predict(model_nb,test.data_nb)
> accuracy(test.data_nb$Purchase,pred)
[1] 0.8362275
```

**From the above Naïve-Bayes model, we can see that the accuracy is 0.8362275.**

From the above built different Naïve-Bayes models, we can achieve a highest accuracy of **0.8362275**.

Hence, we can consider this model as the best among the Naïve-Bayes Classification models built above.

**6. Building KNN Classification Model:**

- a. Clone the data set to a new variable, 'data\_knn' to use for building KNN classification model.

```

> data_knn<-data
> data_knn[1:5,]
  Gender Age Occupation City_Category Stay_In_Current_City_Years Marital_Status
1      F 0-17         10           A                2              0
2      F 0-17         10           A                2              0
3      F 0-17         10           A                2              0
4      F 0-17         10           A                2              0
5      M 55+         16           C                4+              0

  Product_Category_1 Product_Category_2 Product_Category_3 Purchase
1                   3                9.84                12.7    8370
2                   1                 6.00                14.0   15200
3                  12                9.84                12.7   1422
4                  12               14.00                12.7   1057
5                   8                 9.84                12.7   7969
> |

```

- b. Transform the 'Purchase' variable into a dummy variable which has two values, i.e **1** and **0**. Here, **1** means higher purchase [i.e >Mean Value] amount and **0** means lower purchase [i.e <Mean Value] amount made by a customer.

```

> summary(data_knn$Purchase)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   185    5866    8062   9334   12073   23961
> data_knn$Purchase<-ifelse(data_knn$Purchase>9334,1,0)
> data_knn$Purchase<-as.factor(data_knn$Purchase)
> head(data_knn$Purchase)
[1] 0 1 0 0 0 1
Levels: 0 1
> |

```

- c. Create the dummy variables to the categorical variables 'Gender', 'Age', 'Occupation', 'City\_Category', 'Stay\_In\_Current\_City\_Years', 'Marital\_Status'.

```

> library(dummies)
dummies-1.5.6 provided by Decision Patterns

> data_knn<-dummy.data.frame(data_knn,names=c("Gender","Age","Occupation","City_Category","Stay_In_Current_City_Years","Marital_Status"))
> head(data_knn)
  GenderF GenderM Age0-17 Age18-25 Age26-35 Age36-45 Age46-50 Age51-55 Age55+ Occupation0 Occupation1 Occupation2 Occupation3 Occupation4
1      1      0      1      0      0      0      0      0      0      0      0      0      0      0
2      1      0      1      0      0      0      0      0      0      0      0      0      0      0
3      1      0      1      0      0      0      0      0      0      0      0      0      0      0
4      1      0      1      0      0      0      0      0      0      0      0      0      0      0
5      0      1      0      0      0      0      0      0      1      0      0      0      0      0
6      0      1      0      0      0      1      0      0      0      0      0      0      0      0

  Occupation7 Occupation8 Occupation9 Occupation10 Occupation11 Occupation12 Occupation13 Occupation14 Occupation15 Occupation16 Occupation17 Occupation18
1          0          0          0          1          0          0          0          0          0          0          0          0
2          0          0          0          1          0          0          0          0          0          0          0          0
3          0          0          0          1          0          0          0          0          0          0          0          0
4          0          0          0          1          0          0          0          0          0          0          0          0
5          0          0          0          0          0          0          0          0          0          0          1          0
6          0          0          0          0          0          0          0          0          0          1          0          0

  Occupation19 Occupation20 City_CategoryA City_CategoryB City_CategoryC Stay_In_Current_City_Years0 Stay_In_Current_City_Years1 Stay_In_C
1          0          0          1          0          0          0          0          0          0
2          0          0          1          0          0          0          0          0          0
3          0          0          1          0          0          0          0          0          0
4          0          0          1          0          0          0          0          0          0
5          0          0          0          0          1          0          0          0          0

```

- d. Extract numerical variables and normalize the selected data [scaling].



```

> numeric.vars.knn = sapply(data_knn,is.numeric)
> data_knn[numeric.vars.knn] <- lapply(data_knn[numeric.vars.knn],scale)
> head(data_knn)
  GenderF GenderM Age0-17 Age18-25 Age26-35 Age36-45 Age46-50
1 1.7511363 -1.7511363 5.9625827 -0.4710879 -0.8154179 -0.4999519 -0.3005111
2 1.7511363 -1.7511363 5.9625827 -0.4710879 -0.8154179 -0.4999519 -0.3005111
3 1.7511363 -1.7511363 5.9625827 -0.4710879 -0.8154179 -0.4999519 -0.3005111
4 1.7511363 -1.7511363 5.9625827 -0.4710879 -0.8154179 -0.4999519 -0.3005111
5 -0.5710567 0.5710567 -0.1677122 -0.4710879 -0.8154179 -0.4999519 -0.3005111
6 -0.5710567 0.5710567 -0.1677122 -0.4710879 1.2263628 -0.4999519 -0.3005111
  Occupation5 Occupation6 Occupation7 Occupation8 Occupation9 Occupation10 Occu
1 -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025 6.4488003 -(
2 -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025 6.4488003 -(
3 -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025 6.4488003 -(
4 -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025 6.4488003 -(
5 -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025 -0.1550673 -(
6 -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025 -0.1550673 -(
  Occupation18 Occupation19 Occupation20 City_CategoryA City_CategoryB City_Cat
1 -0.1108463 -0.1256246 -0.2553648 1.6482419 -0.8532733 -0.0
2 -0.1108463 -0.1256246 -0.2553648 1.6482419 -0.8532733 -0.0
3 -0.1108463 -0.1256246 -0.2553648 1.6482419 -0.8532733 -0.0
4 -0.1108463 -0.1256246 -0.2553648 1.6482419 -0.8532733 -0.0
5 -0.1108463 -0.1256246 -0.2553648 -0.6067059 -0.8532733 1.
6 -0.1108463 -0.1256246 -0.2553648 1.6482419 -0.8532733 -0.0
  Stay_In_Current_City_Years2 Stay_In_Current_City_Years3 Stay_In_Current_City_
1 2.0988099 -0.4582973 -(
2 2.0988099 -0.4582973 -(
3 2.0988099 -0.4582973 -(
4 2.0988099 -0.4582973 -(
5 -0.4764596 -0.4582973 :
6 -0.4764596 2.1819857 -(
  Product_Category_3 Purchase
1 0.004042377 0
2 0.574222766 1
3 0.004042377 0

```

- e. Dataset being very large (approx. 500 thousand records), we will split data into train data and test data for hold-out evaluation later.

```

> set.seed(123)
> test.index<-1:107515
> train.data_knn<-data_knn[-test.index,]
> test.data_knn<-data_knn[test.index,]
> train.Purchase<-data_knn$Purchase[-test.index]
> test.Purchase<-data_knn$Purchase[test.index]
> head(train.data_knn)
  GenderF GenderM Age0-17 Age18-25 Age26-35
107516 -0.5710567 0.5710567 -0.1677122 -0.4710879 -0.8154179
107517 -0.5710567 0.5710567 -0.1677122 -0.4710879 -0.8154179
107518 -0.5710567 0.5710567 -0.1677122 -0.4710879 1.2263628
107519 -0.5710567 0.5710567 -0.1677122 -0.4710879 1.2263628
107520 -0.5710567 0.5710567 -0.1677122 -0.4710879 1.2263628
107521 -0.5710567 0.5710567 -0.1677122 -0.4710879 1.2263628
  Occupation5 Occupation6 Occupation7 Occupation8 Occupa
107516 -0.151006 -0.1956641 2.8809129 -0.05331976 -0.10
107517 -0.151006 -0.1956641 2.8809129 -0.05331976 -0.10
107518 -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.10
107519 -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.10
107520 -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.10
107521 -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.10
  Occupation17 Occupation18 Occupation19 Occupation20 Ci
107516 -0.2800306 -0.1108463 -0.1256246 -0.2553648
107517 -0.2800306 -0.1108463 -0.1256246 -0.2553648
107518 -0.2800306 -0.1108463 -0.1256246 -0.2553648
107519 -0.2800306 -0.1108463 -0.1256246 -0.2553648
107520 -0.2800306 -0.1108463 -0.1256246 -0.2553648

```

```
> head(test.data_knn)
  GenderF GenderM Age0-17 Age18-25 Age26-35 Age
1  1.7511363 -1.7511363  5.9625827 -0.4710879 -0.8154179 -0.49
2  1.7511363 -1.7511363  5.9625827 -0.4710879 -0.8154179 -0.49
3  1.7511363 -1.7511363  5.9625827 -0.4710879 -0.8154179 -0.49
4  1.7511363 -1.7511363  5.9625827 -0.4710879 -0.8154179 -0.49
5 -0.5710567  0.5710567 -0.1677122 -0.4710879 -0.8154179 -0.49
6 -0.5710567  0.5710567 -0.1677122 -0.4710879  1.2263628 -0.49
  Occupation5 Occupation6 Occupation7 Occupation8 Occupation9
1  -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025
2  -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025
3  -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025
4  -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025
5  -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025
6  -0.151006 -0.1956641 -0.3471115 -0.05331976 -0.1076025
  Occupation18 Occupation19 Occupation20 City_CategoryA City_C
1  -0.1108463 -0.1256246 -0.2553648  1.6482419 -0
2  -0.1108463 -0.1256246 -0.2553648  1.6482419 -0
3  -0.1108463 -0.1256246 -0.2553648  1.6482419 -0
4  -0.1108463 -0.1256246 -0.2553648  1.6482419 -0
5  -0.1108463 -0.1256246 -0.2553648 -0.6067059 -0
6  -0.1108463 -0.1256246 -0.2553648  1.6482419 -0
  Stay_In_Current_City_Years2 Stay_In_Current_City_Years3 Stay
1                2.0988099                -0.4582973
2                2.0988099                -0.4582973
3                2.0988099                -0.4582973
4                2.0988099                -0.4582973
```

- f. Calculate the Accuracy of the KNN Classification models and find the best model.

```
> library(Metrics)
> accuracy(test.Purchase,knn.1)
[1] 0.9862159
> knn.5<-knn(train.data_knn,test.data_knn,train.Purchase,k=5)
> accuracy(test.Purchase,knn.5)
[1] 0.9556248
> knn.15<-knn(train.data_knn,test.data_knn,train.Purchase,k=15)
> accuracy(test.Purchase,knn.15)
[1] 0.910366
```

The largest accuracy of the KNN model is **0.9862** at **K=1**.

## 7. Building Logistic Regression Model

- a. Clone the data set to a new variable, 'data\_lg' to use for building Logistic Regression model.

```
> data_lg<-data
> data_lg[1:5,]
  Gender Age Occupation City_Category Stay_In_Current_City_Years Marital_Status
1     F 0-17         10             A                2              0
2     F 0-17         10             A                2              0
3     F 0-17         10             A                2              0
4     F 0-17         10             A                2              0
5     M 55+         16             C                4+              0
  Product_Category_1 Product_Category_2 Product_Category_3 Purchase
1                 3                9.84                12.7    8370
2                 1                6.00                14.0   15200
3                12                9.84                12.7   1422
4                12               14.00                12.7   1057
5                 8                9.84                12.7   7969
> |
```

- b. Transform the 'Purchase' variable into a dummy variable which has two values, i.e **1** and **0**. Here, **1** means higher purchase [i.e >Mean Value] amount and **0** means lower purchase [i.e <Mean Value] amount made by a customer.

```
> summary(data_lg$Purchase)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   185    5866    8062    9334   12073   23961
> data_lg$Purchase<-ifelse(data_lg$Purchase>9334,1,0)
> data_lg$Purchase<-as.factor(data_lg$Purchase)
> head(data_lg$Purchase)
[1] 0 1 0 0 0 1
Levels: 0 1
> |
```

- c. Split data into train data and test data.

```
> train.index<-createDataPartition(data_lg$Purchase,p=0.8,list=FALSE)
> train.data_lg<-data_lg[train.index,]
> test.data_lg<-data_lg[-train.index,]
> head(train.data_lg)
```

	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
1	F	0-17	10	A	2	0
3	F	0-17	10	A	2	0
4	F	0-17	10	A	2	0
6	M	26-35	15	A	3	0
7	M	46-50	7	B	2	1
8	M	46-50	7	B	2	1

```

  Product_Category_1 Product_Category_2 Product_Category_3 Purchase
1                   3                   9.84                12.7      0
3                   12                   9.84                12.7      0
4                   12                  14.00                12.7      0
6                    1                    2.00                12.7      1
7                    1                    8.00                17.0      1
8                    1                   15.00                12.7      1
> head(test.data_lg)
```

	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
2	F	0-17	10	A	2	0
5	M	55+	16	C	4+	0
14	M	26-35	20	A	1	1
22	M	26-35	12	C	4+	1
26	M	26-35	17	C	0	0
30	F	36-45	1	B	4+	1

```

  Product_Category_1 Product_Category_2 Product_Category_3 Purchase
2                   1                    6.00                14.0      1
5                    8                    9.84                12.7      0
14                   1                    2.00                 5.0      1
22                   8                    9.84                12.7      1
26                   6                    8.00                12.7      1
30                   2                    4.00                 8.0      1
> |
```

- d. Build full logistic model using 'glm()'.

```

> model_lg_full<-glm(as.factor(Purchase)~.,data=train.data_lg,family=binomial())
> summary(model_lg_full)

Call:
glm(formula = as.factor(Purchase) ~ ., family = binomial(), data = train.data_lg)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.6622  -0.9797  -0.7115   1.0496   2.4619

Coefficients:
                Estimate Std. Error z value Pr(>|z|)
(Intercept)    0.3255572   0.0291102   11.184 < 2e-16 ***
GenderM        0.2224393   0.0078314   28.404 < 2e-16 ***
Age18-25       0.1153657   0.0215789    5.346 8.98e-08 ***
Age26-35       0.1945145   0.0209692    9.276 < 2e-16 ***
Age36-45       0.2572174   0.0215513   11.935 < 2e-16 ***
Age46-50       0.2499432   0.0236833   10.554 < 2e-16 ***
Age51-55       0.3814411   0.0241499   15.795 < 2e-16 ***
Age55+         0.3456951   0.0264662   13.062 < 2e-16 ***
Occupation     0.0023126   0.0005136    4.502 6.72e-06 ***
City_CategoryB 0.0651329   0.0082552    7.890 3.02e-15 ***
City_CategoryC 0.2675380   0.0088579   30.203 < 2e-16 ***
Stay_In_Current_City_Years1 0.0159197   0.0106084    1.501 0.133440
Stay_In_Current_City_Years2 0.0184713   0.0118245    1.562 0.118262
Stay_In_Current_City_Years3 -0.0022555   0.0120262   -0.188 0.851233
Stay_In_Current_City_Years4+ 0.0322954   0.0123173    2.622 0.008743 **
Marital_Status -0.0245683   0.0071527   -3.435 0.000593 ***
Product_Category_1 -0.1728760   0.0011102 -155.722 < 2e-16 ***
Product_Category_2 -0.0486791   0.0008902  -54.686 < 2e-16 ***
Product_Category_3  0.0073419   0.0014769    4.971 6.65e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 579321  on 430061  degrees of freedom
Residual deviance: 529299  on 430043  degrees of freedom
AIC: 529337

Number of Fisher Scoring iterations: 4

> |

```

- e. Build a base model with one x-variable.

```
> model_lg_base<-glm(as.factor(Purchase)~Gender,data=train.data_lg,family=binomial())
> summary(model_lg_base)
```

Call:

```
glm(formula = as.factor(Purchase) ~ Gender, family = binomial(),
    data = train.data_lg)
```

Deviance Residuals:

```
      Min       1Q   Median       3Q      Max
-1.0405  -1.0405  -0.9281   1.3208   1.4491
```

Coefficients:

```
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.619282    0.006449  -96.03   <2e-16 ***
GenderM      0.288398    0.007366   39.15   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 579321  on 430061  degrees of freedom
Residual deviance: 577765  on 430060  degrees of freedom
AIC: 577769
```

Number of Fisher Scoring iterations: 4

f. Build a stepwise forward model.

```
> model_lg_fwd<-step(model_lg_base,scope=list(upper=model_lg_full,lower=~1),direction="forward",trace=T)
Start:  AIC=577769.4
as.factor(Purchase) ~ Gender
```

	Df	Deviance	AIC
+ Product_Category_1	1	534227	534233
+ Product_Category_2	1	559375	559381
+ City_Category	2	576216	576224
+ Product_Category_3	1	576723	576729
+ Age	6	577612	577628
+ Occupation	1	577699	577705
+ Stay_In_Current_City_Years	4	577739	577751
<none>		577765	577769
+ Marital_Status	1	577765	577771

Step: AIC=534232.8

```
as.factor(Purchase) ~ Gender + Product_Category_1
```

	Df	Deviance	AIC
+ Product_Category_2	1	531142	531150
+ City_Category	2	533015	533025
+ Age	6	533594	533612
+ Product_Category_3	1	534066	534074
+ Occupation	1	534169	534177
+ Marital_Status	1	534205	534213
+ Stay_In_Current_City_Years	4	534203	534217
<none>		534227	534233

Step: AIC=531149.6

```
as.factor(Purchase) ~ Gender + Product_Category_1 + Product_Category_2
```

	Df	Deviance	AIC
+ City_Category	2	529914	529926
+ Age	6	530434	530454
+ Occupation	1	531080	531090
+ Product_Category_3	1	531112	531122
+ Marital_Status	1	531117	531127
+ Stay_In_Current_City_Years	4	531116	531132

<none> 531142 531150

Step: AIC=529926.4

as.factor(Purchase) ~ Gender + Product\_Category\_1 + Product\_Category\_2 +  
City\_Category

	Df	Deviance	AIC
+ Age	6	529367	529391
+ Occupation	1	529871	529885
+ Product_Category_3	1	529885	529899
+ Marital_Status	1	529901	529915
+ Stay_In_Current_City_Years	4	529899	529919
<none>		529914	529926

Step: AIC=529391.4

as.factor(Purchase) ~ Gender + Product\_Category\_1 + Product\_Category\_2 +  
City\_Category + Age

	Df	Deviance	AIC
+ Product_Category_3	1	529343	529369
+ Occupation	1	529347	529373
+ Marital_Status	1	529356	529382
+ Stay_In_Current_City_Years	4	529355	529387
<none>		529367	529391

Step: AIC=529368.6

as.factor(Purchase) ~ Gender + Product\_Category\_1 + Product\_Category\_2 +  
City\_Category + Age + Product\_Category\_3

	Df	Deviance	AIC
+ Occupation	1	529322	529350
+ Marital_Status	1	529331	529359
+ Stay_In_Current_City_Years	4	529331	529365
<none>		529343	529369

Step: AIC=529350.1

as.factor(Purchase) ~ Gender + Product\_Category\_1 + Product\_Category\_2 +  
City\_Category + Age + Product\_Category\_3 + Occupation

	Df	Deviance	AIC
+ Marital_Status	1	529310	529340
+ Stay_In_Current_City_Years	4	529310	529346
<none>		529322	529350

Step: AIC=529340.3

as.factor(Purchase) ~ Gender + Product\_Category\_1 + Product\_Category\_2 +  
City\_Category + Age + Product\_Category\_3 + Occupation + Marital\_Status

	Df	Deviance	AIC
+ Stay_In_Current_City_Years	4	529299	529337
<none>		529310	529340

Step: AIC=529336.5

as.factor(Purchase) ~ Gender + Product\_Category\_1 + Product\_Category\_2 +  
City\_Category + Age + Product\_Category\_3 + Occupation + Marital\_Status +  
Stay\_In\_Current\_City\_Years

```

> summary(model_lg_fwd)

Call:
glm(formula = as.factor(Purchase) ~ Gender + Product_Category_1 +
  Product_Category_2 + City_Category + Age + Product_Category_3 +
  Occupation + Marital_Status + Stay_In_Current_City_Years,
  family = binomial(), data = train.data_lg)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.6622  -0.9797  -0.7115   1.0496   2.4619

Coefficients:
                Estimate Std. Error z value Pr(>|z|)
(Intercept)      0.3255572   0.0291102   11.184 < 2e-16 ***
GenderM          0.2224393   0.0078314   28.404 < 2e-16 ***
Product_Category_1 -0.1728760   0.0011102 -155.722 < 2e-16 ***
Product_Category_2 -0.0486791   0.0008902  -54.686 < 2e-16 ***
City_CategoryB    0.0651329   0.0082552    7.890 3.02e-15 ***
City_CategoryC    0.2675380   0.0088579   30.203 < 2e-16 ***
Age18-25         0.1153657   0.0215789    5.346 8.98e-08 ***
Age26-35         0.1945145   0.0209692    9.276 < 2e-16 ***
Age36-45         0.2572174   0.0215513   11.935 < 2e-16 ***
Age46-50         0.2499432   0.0236833   10.554 < 2e-16 ***
Age51-55         0.3814411   0.0241499   15.795 < 2e-16 ***
Age55+          0.3456951   0.0264662   13.062 < 2e-16 ***
Product_Category_3 0.0073419   0.0014769    4.971 6.65e-07 ***
Occupation       0.0023126   0.0005136    4.502 6.72e-06 ***
Marital_Status   -0.0245683   0.0071527   -3.435 0.000593 ***
Stay_In_Current_City_Years1 0.0159197   0.0106084    1.501 0.133440
Stay_In_Current_City_Years2 0.0184713   0.0118245    1.562 0.118262
Stay_In_Current_City_Years3 -0.0022555   0.0120262   -0.188 0.851233
Stay_In_Current_City_Years4+ 0.0322954   0.0123173    2.622 0.008743 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 579321  on 430061  degrees of freedom
Residual deviance: 529299  on 430043  degrees of freedom
AIC: 529337

Number of Fisher Scoring iterations: 4

> |

```

- g. Build a stepwise backward model using the full model built previously.



```
> model_lg_bwd<-step(model_lg_full,direction="backward",trace=T)
Start: AIC=529336.5
as.factor(Purchase) ~ Gender + Age + Occupation + City_Category +
  Stay_In_Current_City_Years + Marital_Status + Product_Category_1 +
  Product_Category_2 + Product_Category_3
```

	Df	Deviance	AIC
<none>		529299	529337
- Stay_In_Current_City_Years	4	529310	529340
- Marital_Status	1	529310	529346
- Occupation	1	529319	529355
- Product_Category_3	1	529323	529359
- Age	6	529816	529842
- Gender	1	530112	530148
- City_Category	2	530352	530386
- Product_Category_2	1	532321	532357
- Product_Category_1	1	557414	557450

```
> summary(model_lg_bwd)
```

Call:

```
glm(formula = as.factor(Purchase) ~ Gender + Age + Occupation +
  City_Category + Stay_In_Current_City_Years + Marital_Status +
  Product_Category_1 + Product_Category_2 + Product_Category_3,
  family = binomial(), data = train.data_lg)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.6622	-0.9797	-0.7115	1.0496	2.4619

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.3255572	0.0291102	11.184	< 2e-16 ***
GenderM	0.2224393	0.0078314	28.404	< 2e-16 ***
Age18-25	0.1153657	0.0215789	5.346	8.98e-08 ***
Age26-35	0.1945145	0.0209692	9.276	< 2e-16 ***
Age36-45	0.2572174	0.0215513	11.935	< 2e-16 ***
Age46-50	0.2499432	0.0236833	10.554	< 2e-16 ***
Age51-55	0.3814411	0.0241499	15.795	< 2e-16 ***
Age55+	0.3456951	0.0264662	13.062	< 2e-16 ***
Occupation	0.0023126	0.0005136	4.502	6.72e-06 ***
City_CategoryB	0.0651329	0.0082552	7.890	3.02e-15 ***
City_CategoryC	0.2675380	0.0088579	30.203	< 2e-16 ***
Stay_In_Current_City_Years1	0.0159197	0.0106084	1.501	0.133440
Stay_In_Current_City_Years2	0.0184713	0.0118245	1.562	0.118262
Stay_In_Current_City_Years3	-0.0022555	0.0120262	-0.188	0.851233
Stay_In_Current_City_Years4+	0.0322954	0.0123173	2.622	0.008743 **
Marital_Status	-0.0245683	0.0071527	-3.435	0.000593 ***
Product_Category_1	-0.1728760	0.0011102	-155.722	< 2e-16 ***
Product_Category_2	-0.0486791	0.0008902	-54.686	< 2e-16 ***
Product_Category_3	0.0073419	0.0014769	4.971	6.65e-07 ***

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 579321 on 430061 degrees of freedom
Residual deviance: 529299 on 430043 degrees of freedom
AIC: 529337
```

Number of Fisher Scoring iterations: 4




- h. Use AIC as metric to conclude a better model.

Using AIC as metric: From the above build stepwise forward and backward models, we can see that both the approaches suggest the same model, which is similar to full model.

- i. Calculate the Accuracy of the final Logistic Regression model.

```
> model_lg_fit<-train(as.factor(Purchase)~.,data=train.data_lg,method="glm",family="binomial")
> pred<-predict(model_lg_fit,newdata=test.data_lg)
> accuracy(pred,test.data_lg$Purchase)
[1] 0.7606567
> |
```

## 8. Area Under curve for Best Model



```
> prediction_val_nb = prediction(as.numeric(pred),as.numeric(test.data_nb$Purchase))
> performance_nb = performance(prediction_val_nb,"tpr","fpr")
> plot(performance_nb,col="blue")
```

```
> knn.1<-knn(train.data_knn,test.data_knn,train.Purchase,k=1)
> prediction_val_knn <- prediction(as.numeric(knn.1),as.numeric(test.data_knn$Purchase))
> performance_knn = performance(prediction_val_knn,"tpr","fpr")
> plot(performance_nb,col = "red")
> plot(performance_knn,add = TRUE,col = "blue")
> |
```

```
> AUC_knn = max(attr(performance_knn,"y.values")[[1]]-(attr(performance_knn,"x.values")[[1]]))
> AUC_knn
[1] 0.9715863
> AUC_nb = max(attr(performance_nb,"y.values")[[1]]-(attr(performance_nb,"x.values")[[1]]))
> AUC_nb
[1] 0.6582353
> |
```